



Stage A · Encoder — Spatio-temporal representation learning · 33,026 weights



Stage B · Decoder + Loss — Mixture RUM choice probability + cross-entropy training · 77 interp. params



Stage C · Inference — Counterfactual deployment + BPR closure · Θ frozen, no re-fit

