


Figure 2 — ABM Forward Simulation Pipeline

agent-based deployment with destination-side BPR congestion closure



Stage B instantiates N heterogeneous agents, each fetching their personal $(\beta_{t, m_n}, \kappa_{k_n}, \delta_{k_n})$ from frozen Θ and choosing one destination via Gumbel-max. Stage C then closes the destination-side congestion loop (BPR multiplier on car) iteratively via MSA until rel-gap $< 0.5\%$. The same trained Θ from Figure 1 is used; Stage C does no re-fitting.